

The Art of Relevance and Recommendations

@cchio

slides+code: <https://git.io/vXRwv>

**...i.e. how to hit the ground running
with recommender systems**

What do I do?

- **Stanford B.S./M.S. Computer Science**
- **Research Engineer at Shape Security**
- **Organizer, 'Data Mining for Cyber Security' meetup group in Silicon Valley**
- **Authoring 'Machine Learning and Security' (O'Reilly, mid-late 2017)**
- **Presented independent research in Adversarial ML at dozens of security/ML conferences over last two years**



- Bought Justin Bieber
- Bought Selena Gomez



- Performs a search for Justin Bieber
- Recommender system suggests Selena Gomez

facebook®

amazon

 Spotify®

last.fm

You



 yelp

Google

NETFLIX



Linked



twitter

ebay

bing™

 stackoverflow



reddit



airbnb

Ways to serve recommendations (in desc. order of naivety)

1. Hand-crafted; “our picks”
2. Global aggregates; top 10 among all users, most recent items
3. Individualized recommendations based on inferred/learned user preferences

MovieLens dataset

- 1. Activity from MovieLens, a movie recommendation service**
- 2. 100004 ratings, 1296 tags, 9125 movies**
- 3. 671 users, all selected users rated at least 20 movies**

How to guess what a user might like?

Explicit

★★★★★ Very sharp !!! Excellent for cheese!

By [Kyle & Rachael](#) on March 20, 2016

Verified Purchase

Works alot better than I had expected, I previously had a kitchen aid zester that I purchased this one. I'm very impressed with the quality and sharpness of it, works great. Washes very easily and comes with a convenient cover. Tested it out with zesting cheese. Passed with flying colors.

Yes id recommend this product to anyone

Implicit

Customers Who Bought This Item Also Bought



Zulay Premium Quality
Metal Lemon Lime
Squeezer - Manual Citrus
Juice Press

★★★★★ 230

\$12.95 ✓ Prime


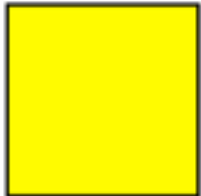


Deiss PRO Dual Julienne
Peeler & Vegetable Peeler
— Non-slip Comfortable
Handle — Amazing Tool...

★★★★★ 382

\$9.88 ✓ Prime

How to guess what a user might like?

		users											
		1	2	3	4	5	6	7	8	9	10	11	12
movies	1	1		3			5			5		4	
	2			5	4			4			2	1	3
	3	2	4		1	2		3		4	3	5	
	4		2	4		5			4			2	
	5			4	3	4	2					2	5
	6	1		3		3			2			4	
		 - unknown rating		 - rating between 1 to 5									

Real Problems

1. **Cold-start** for new users/items
2. Not leveraging on **crowd wisdom**
3. **Sparse** ratings
4. **Popularity bias**/Long-tail problem
5. Computational **efficiency** issues

Types of Recommender Systems


1. Content based filtering
2. Collaborative filtering
3. Latent-factor based
4. Multi-armed bandits

Content based filtering

1. Recommendations made based solely on the characteristics of the item
2. Build a **profile of a user based on his past likes**
3. System recommends items similar to what a user has liked in the past


Content based filtering

solves



1. New items can be recommended
2. Obscure items will not be neglected
3. Can be efficiently pre-computed

but



- Finding good features is hard/laborious
- Cannot recommend to new users without a profile

Collaborative filtering

User-based

1. Get the active user's item ratings (explicit + implicit)
2. Identify other users that made similar ratings on similar items
3. Recognize items that these similar users liked
4. Generate a prediction for the active user


Collaborative filtering

Item-based

1. Look at items the active user rated
2. Compute the similarity of these items to other items, based on how other users rated all of them
3. Generate a prediction for the active user


Collaborative filtering

solves



1. Leverages heavily on crowd wisdom
2. No feature engineering required

but




- Cold-start
- Sparse ratings
- Cannot recommend obscure items
- Hard to pre-compute, ratings constantly being updated

Latent-factor based

- 1. Develop a model of user-item relations that tries to uncover latent/hidden relationships or preferences**
- 2. Use ML algorithms to model similarities**

Latent-factor based

Low-rank matrix factorization



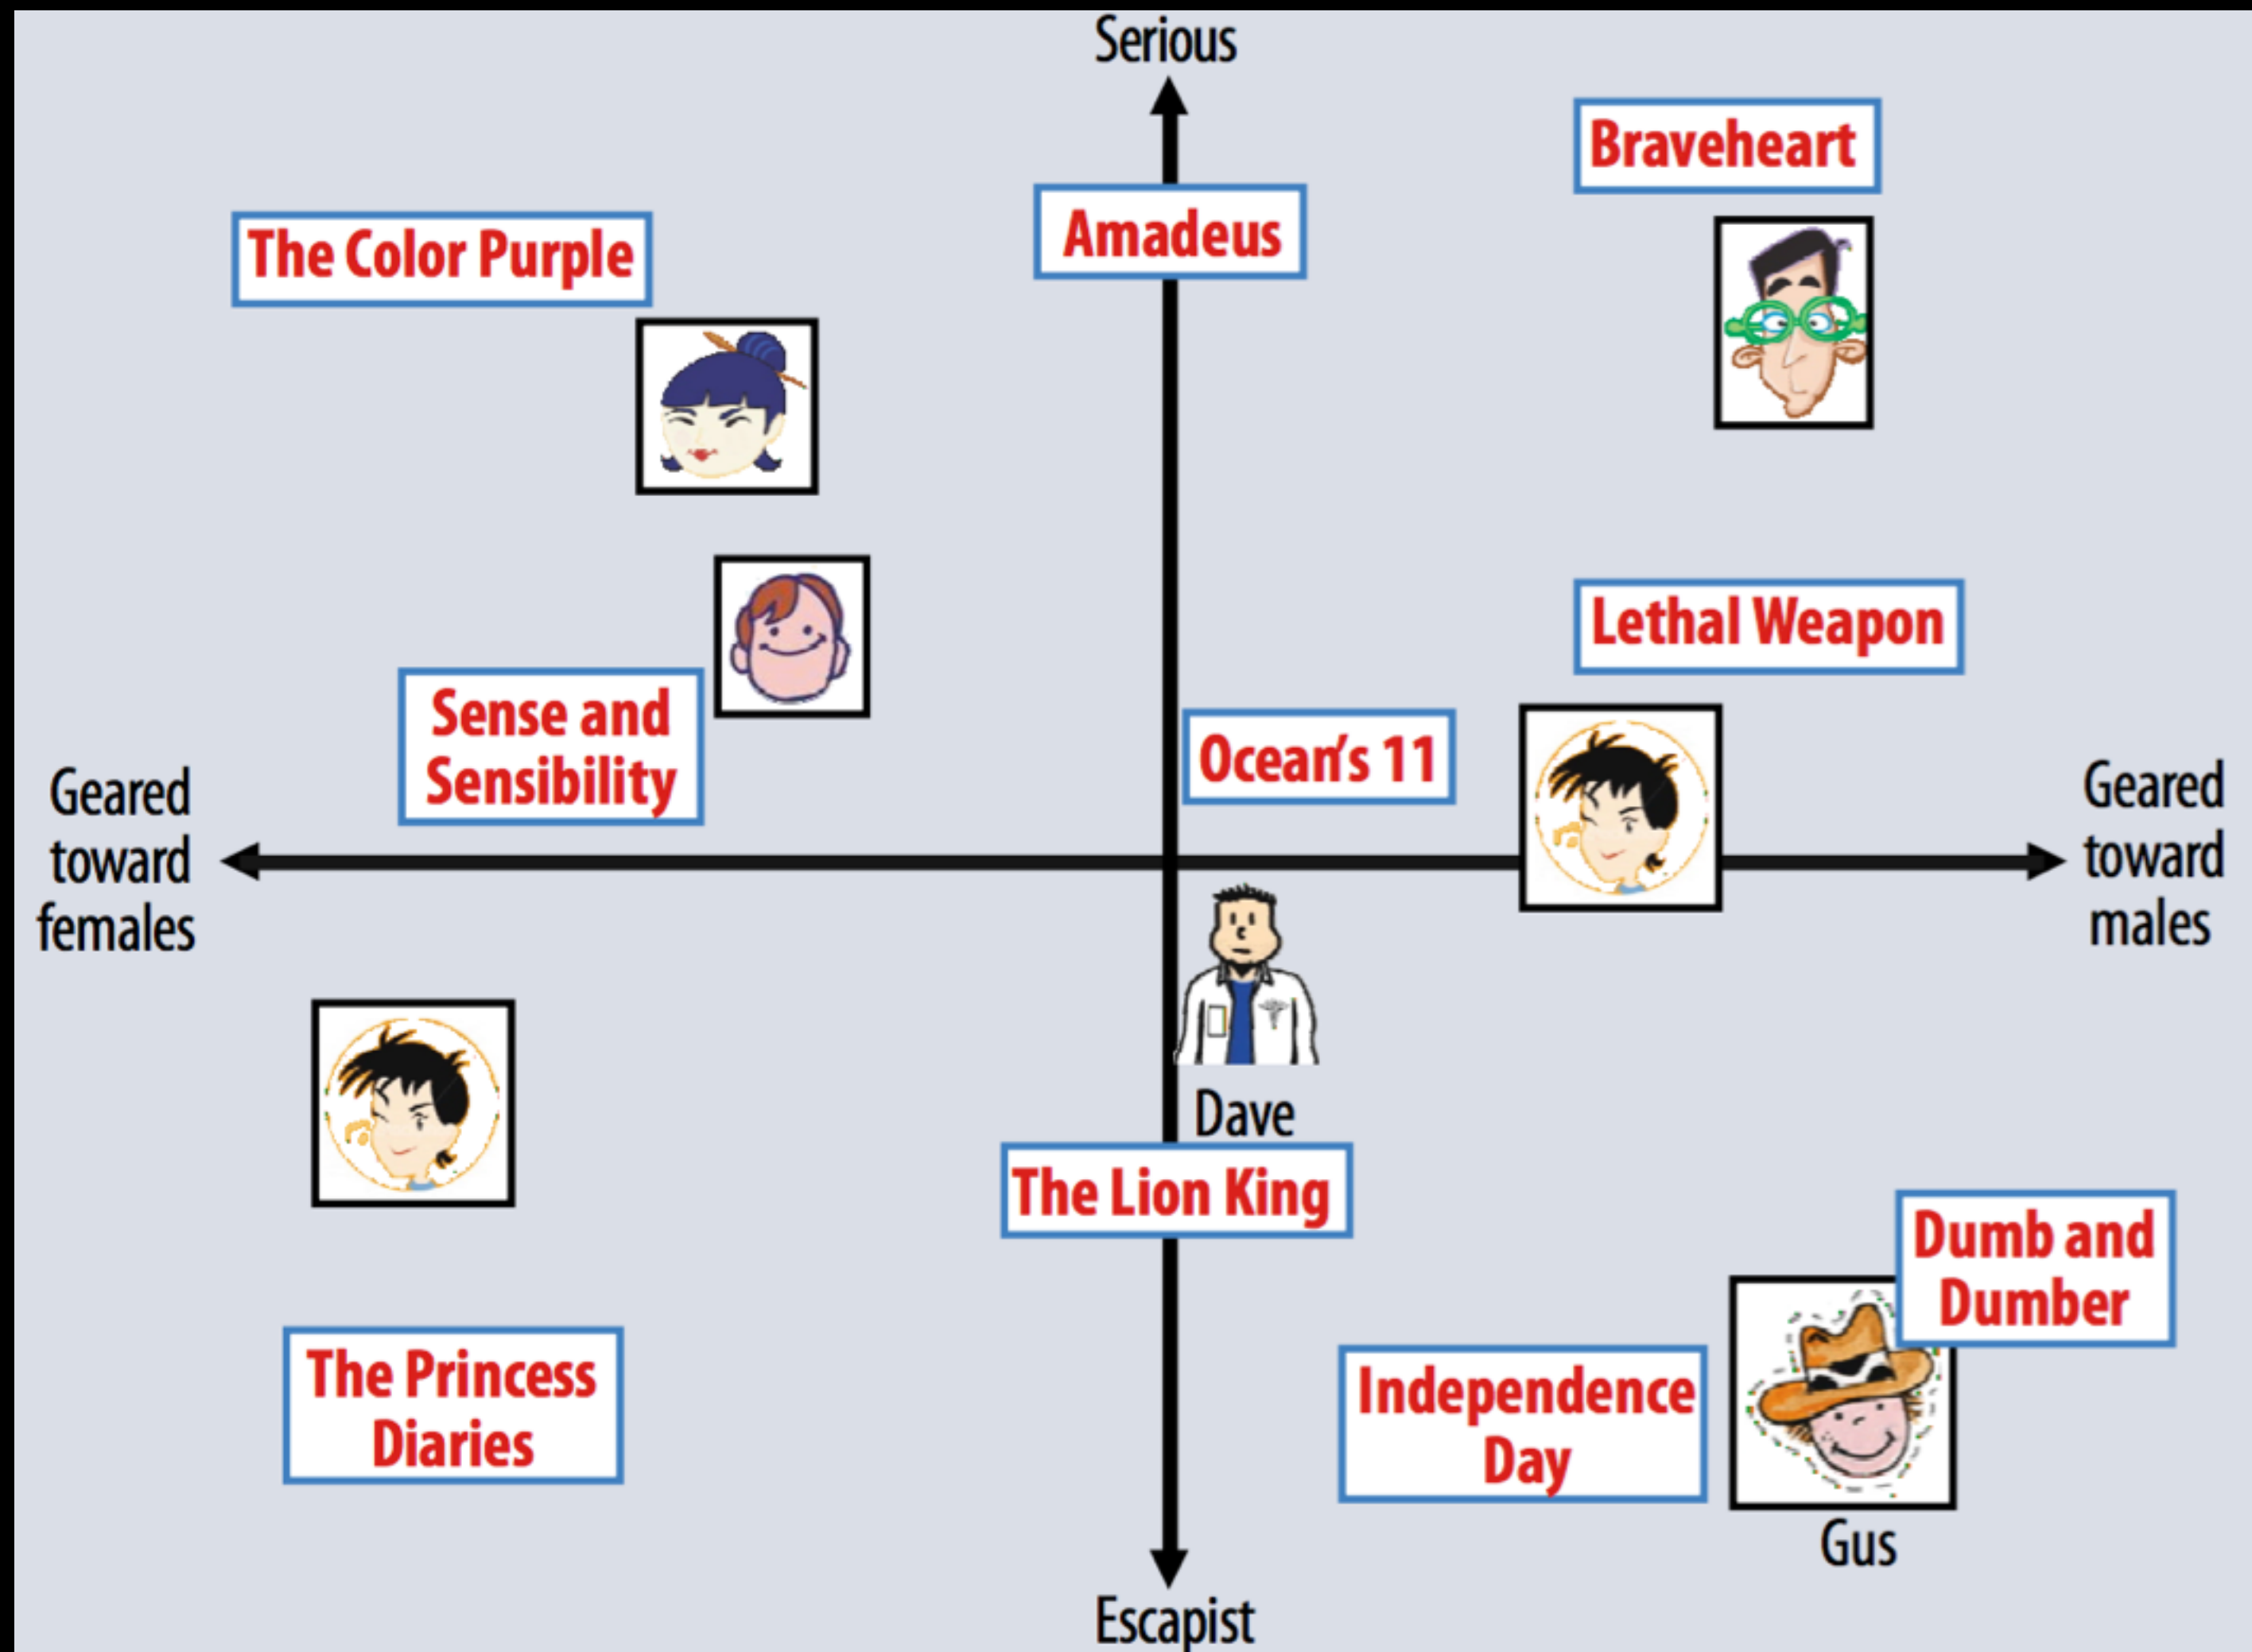



Figure 2. A simplified illustration of the latent factor approach, which characterizes both users and movies using two axes—male versus female and serious versus escapist.


Latent-factor based

solves




1. Fixes sparse ratings
2. Similarity comparisons more efficient

but



- Still not enough recommending new/
unique items to users (long-tail)


The multi-armed bandit problem




1. Each arm x
 - A. Wins (output = 1) with a fixed unknown probability μ_x
 - B. Loses (output = 0) with a fixed unknown probability $1-\mu_x$
2. All you have is posterior probability
3. How to pull arms to **maximize total reward?**



μ_1



μ_2



...

...

...



μ_k



Naive solution

1. Pull each arm k times (e.g. 100)
2. Record the mean reward of each arm
3. Pull the arm with the highest mean reward for eternity

Thompson Sampling works real well

1. Maintain a **probability function for each arm, based on data collected over time**
2. Draw a sample from every arm's probability function
3. Pull the arm that gives the largest drawn sample
4. Repeat for eternity

How does this apply to **content A/B testing**?

k arms → **k unique content versions**

goal → **maximize user retention, CTR**

How does this apply to **online advertising**?

k arms → **k distinct page visitors**
goal → **maximize ad engagement**

How does this apply to **clinical trials**?

k arms → **k distinct treatments**

goal → **minimize patient loss**

How does this apply to **recommender systems**?

k arms → **k distinct items to recommend**

goal → **maximize user interest**

...you're running an experiment each time you pull an arm

How does this apply to **recommender systems**?

k arms → **k distinct items to recommend**

goal → **maximize user interest**

...you're running an experiment each time you pull an arm

Evaluation

Practical evaluation concerns

$$\sqrt{\sum_{xi} (r_{xi} - r_{xi}^*)^2}$$

RMSE unreasonably penalizes a method that **does well for predicting high ratings** but not low ratings

vs.

- Precision @ Top-5
- Receiver Operating Characteristic (ROC) Area Under Curve (AUC)

Practical evaluation concerns

**Prediction diversity, context, order
are often as important as accuracy**

Practical evaluation concerns

Cost function - what happens when you make a bad recommendation?

Complexity/Efficiency

- Expensive step is **item/customer similarity matching**
- Often too expensive to do at runtime
 - Precompute, good for item-item collab filtering
- Dimensionality reduction
- Locality sensitive hashing
- Clustering

Questions?

Takeaways

1. Implementing collaborative filtering is much easier than you think
2. Use matrix factorization to fix sparse ratings
3. Exploitation vs. Exploration
→ using multi-armed bandit algorithms

Practical tips

1. **Implicit signals > explicit signals**
2. **Optimize for model explainability instead of just accuracy**
3. **(More data + simple algorithm) > (less data + complex algorithm)**

slides+code: <https://git.io/vXRwv>
@cchio